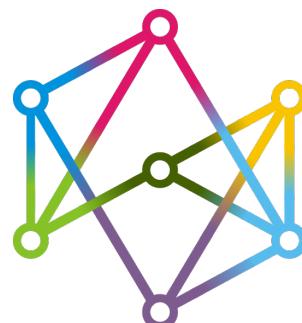


Online Simplified Belief Space Planning with Performance Guarantees

Vadim Indelman



Technion Robotics Expo

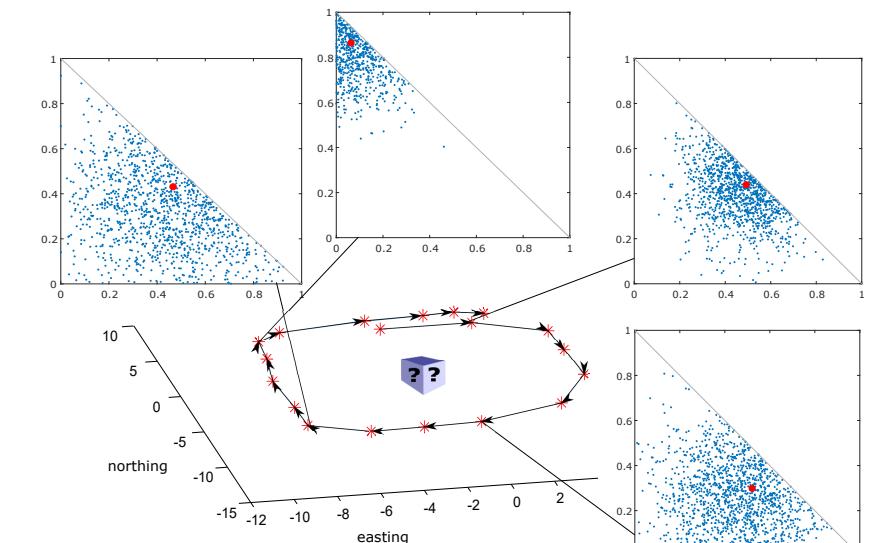
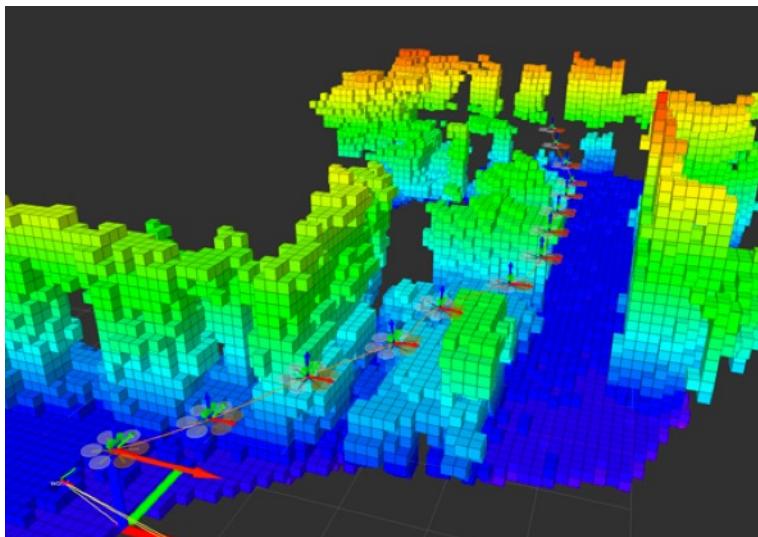
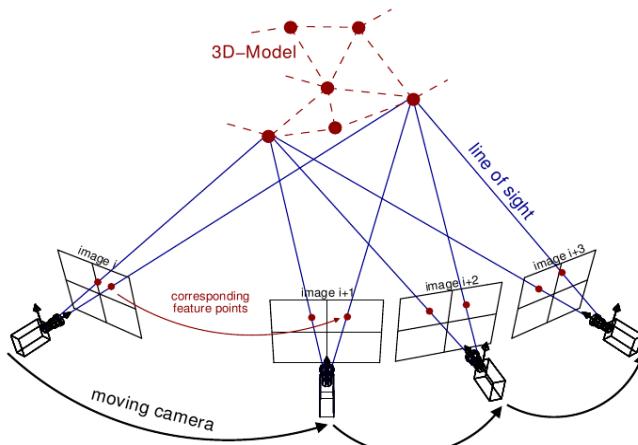


June 12, 2022

Introduction

Key required capabilities:

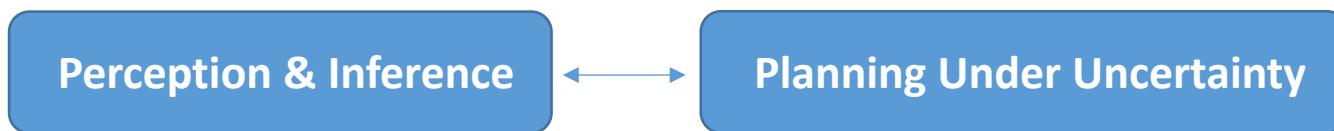
- **Perception and Inference:** Where am I? What is the surrounding environment?



Introduction

Key required capabilities:

- **Perception and Inference:** Where am I? What is the surrounding environment?
- **Planning Under Uncertainty:** What should I be doing next?
 - Determine best action(s) to accomplish a task, account for different sources of uncertainty



- Related problems: autonomous navigation, active SLAM, informative planning/sensing, etc.



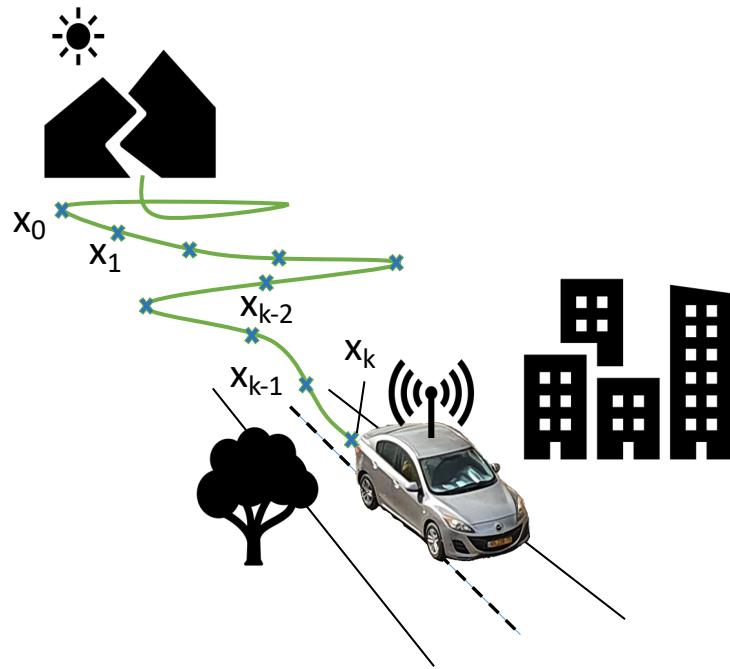
Perception and Inference

- Posterior belief at time k:

$$b[X_k] \doteq \mathbb{P}(X_k \mid u_{0:k-1}, z_{1:k})$$

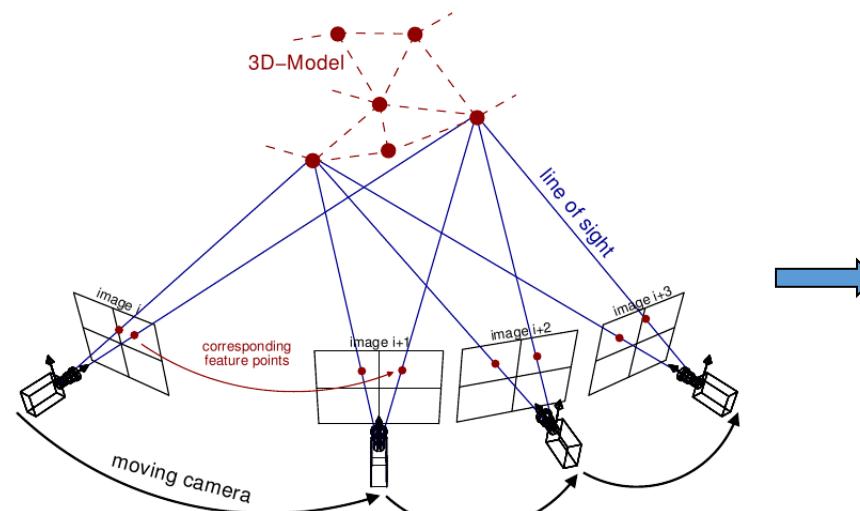
↑
state/variables at
time instant k ↑
user controls ↑
observations

- Example:

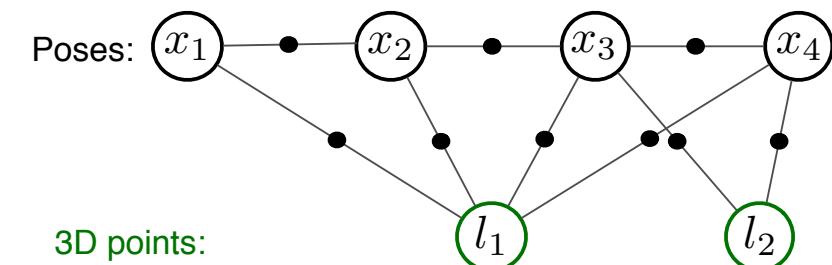


$$X_k \doteq \underline{\{x_0, \dots, x_k\}} \quad \overline{L_k}$$

Past & current
robot states Landmarks



Can be represented with
graphical models, e.g. a Factor Graph



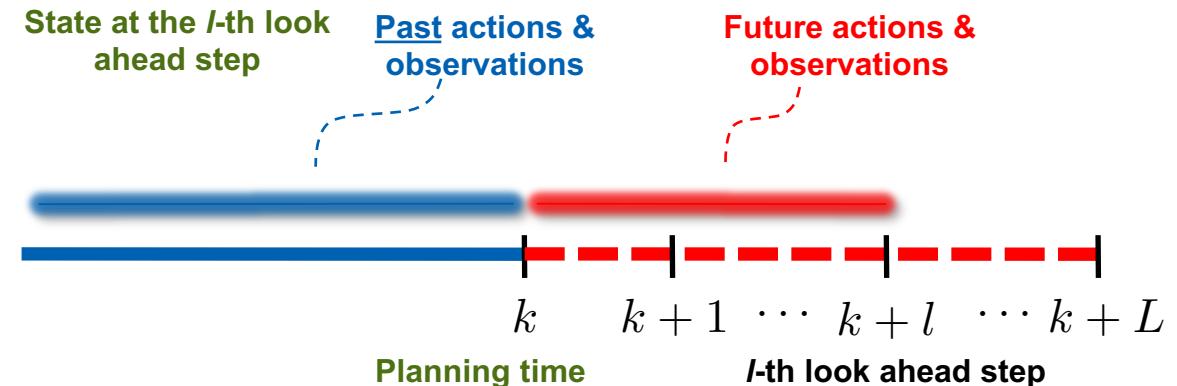
Belief Space Planning (BSP)

- Objective function:

$$J(b[X_k], a) = \mathbb{E}\left\{\sum_{l=0}^{L-1} r(b[X_{k+l}], a_{k+l}) + r(b[X_{k+L}])\right\}, \quad a \doteq a_{k:k+L-1} \in \mathcal{A}$$

- Belief at the l -th look-ahead step:

$$b[X_{k+l}] \doteq \mathbb{P}(X_{k+l} \mid u_{0:k-1}, z_{0:k}, u_{k:k+l-1}, z_{k+1:k+l})$$



- Examples for reward function $r(b, a)$:

- Distance to goal (navigate to a goal)
- Information theoretic reward (reduce uncertainty, exploration)

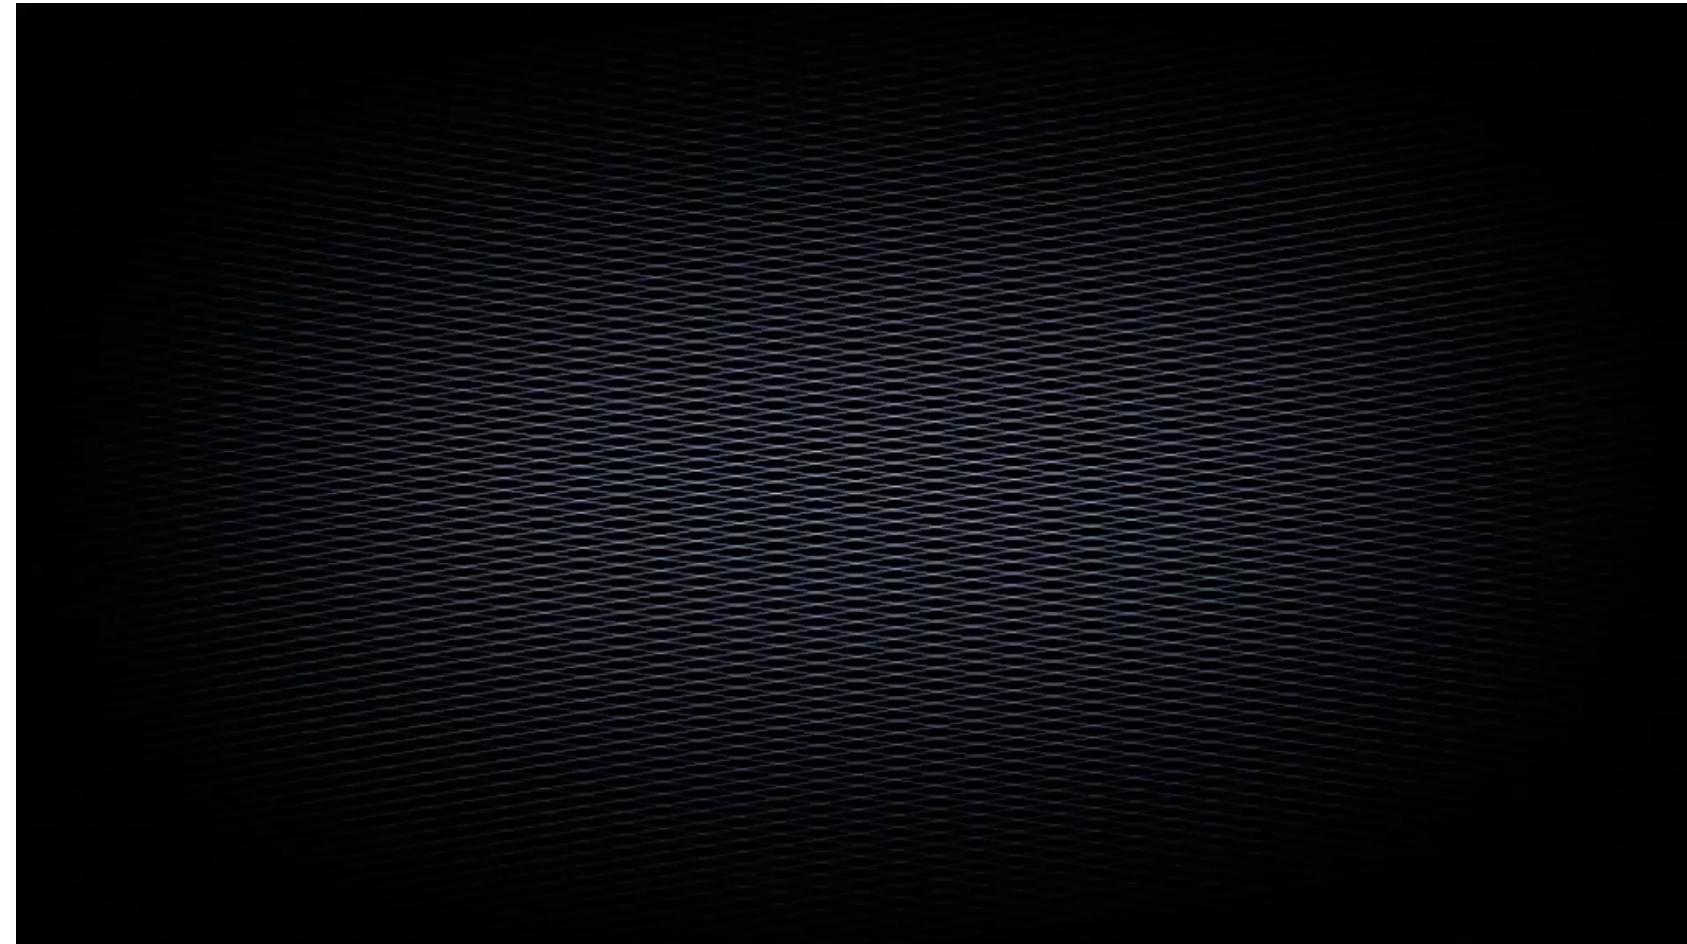
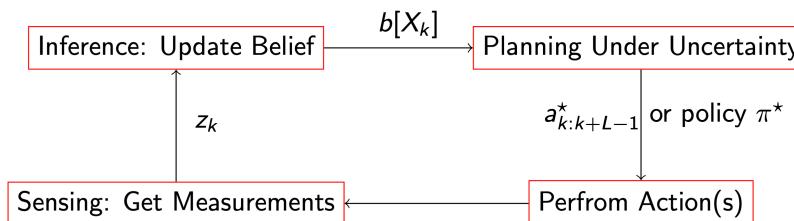


ANPL

Autonomous Navigation
and Perception Lab

Collaborative BSP in Unknown Environments on Quads

Plan-act-sense-infer framework



Belief Space Planning (BSP)

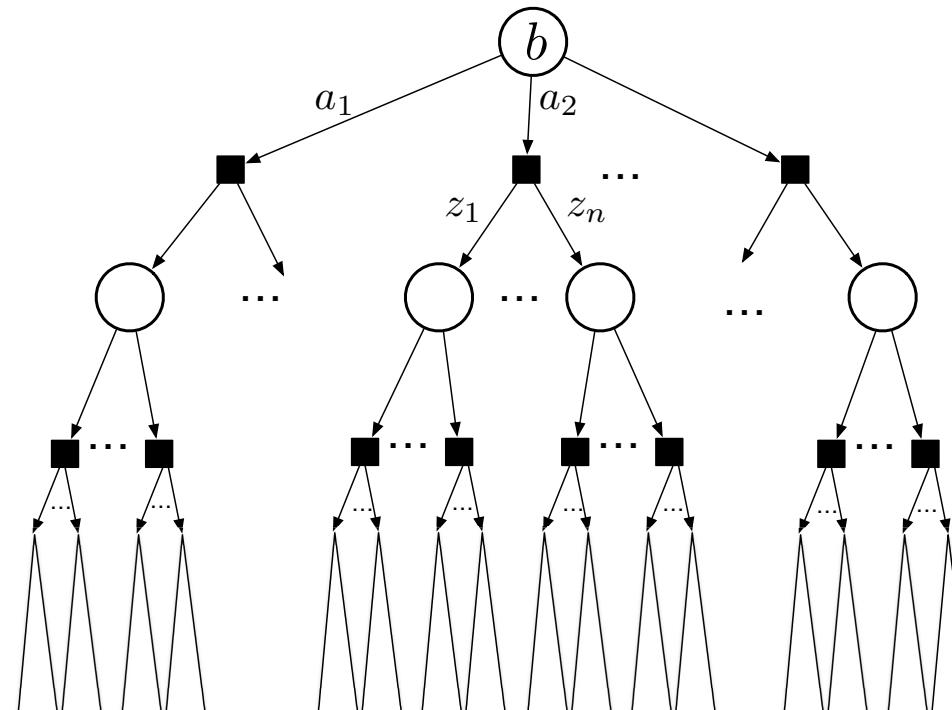
Action sequences:

$$J(b[X_k], a) = \mathbb{E}\left\{\sum_{l=0}^{L-1} r(b[X_{k+l}], a_{k+l}) + r(b[X_{k+L}])\right\}, \quad a \doteq a_{k:k+L-1} \in \mathcal{A}$$

Policies:

$$v_k^\pi(b_k) \equiv J_k(b_k, \pi) = \mathbb{E}\left\{\sum_{l=0}^{L-1} r(b_{k+l}, \pi_{k+l}(b_{k+l})) + r(b_{k+L})\right\}$$

- Finding an optimal solution is generally **computationally intractable**



Belief Space Planning (BSP)

Action sequences:

$$J(b[X_k], a) = \mathbb{E}\left\{\sum_{l=0}^{L-1} r(b[X_{k+l}], a_{k+l}) + r(b[X_{k+L}])\right\}, \quad a \doteq a_{k:k+L-1} \in \mathcal{A}$$

Policies:

$$v_k^\pi(b_k) \equiv J_k(b_k, \pi) = \mathbb{E}\left\{\sum_{l=0}^{L-1} r(b_{k+l}, \pi_{k+l}(b_{k+l})) + r(b_{k+L})\right\}$$

- Finding an optimal solution is generally **computationally intractable**
- Our focus - want to:
 - Act autonomously, online, while accounting for different sources of uncertainty and ambiguity
 - Reliably operate in uncertain, perceptually aliased environments/scenarios
 - Acquire high-level / semantic environment understanding
 - Support **belief-dependent** rewards (e.g. information-theoretic)

Simplification of Decision-Making Problems

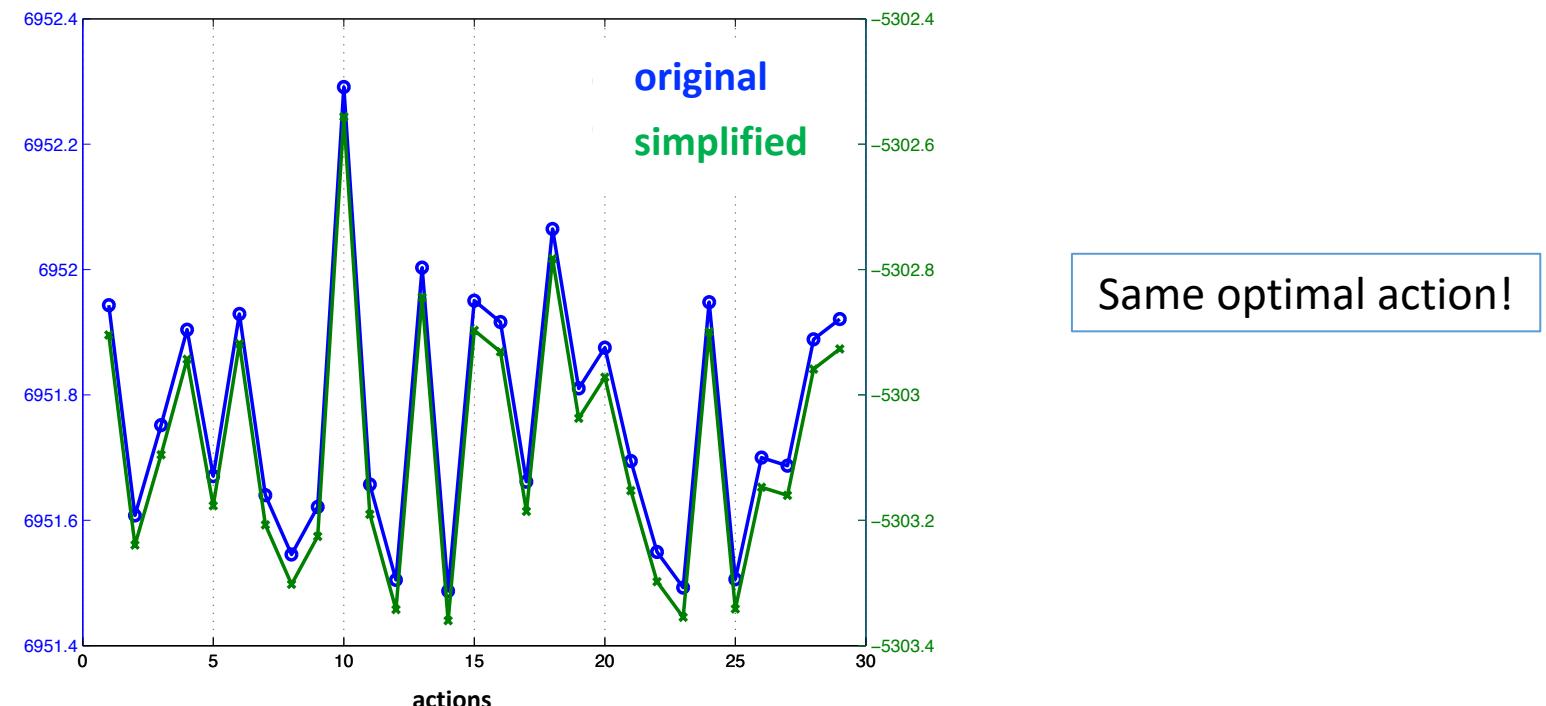
Concept:

- Identify and solve a **simplified computationally-easier** decision-making problem
- Provide performance guarantees – how can we measure simplification quality?

Simplification of Decision-Making Problems

[Indelman RA-L'16][Elimelech and Indelman, IJRR'22]

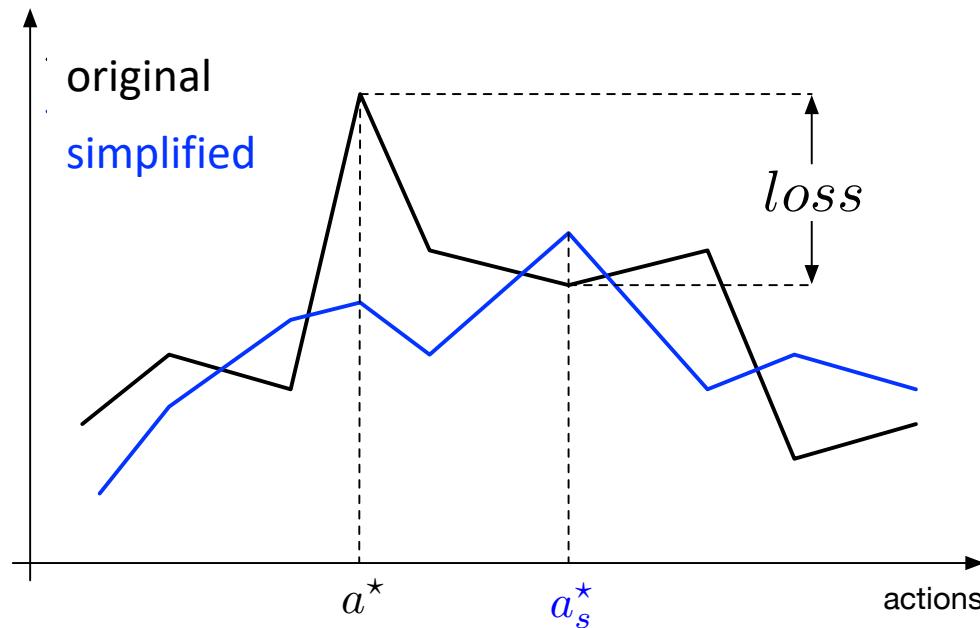
- **Key observation:** In decision making, only need to sort actions from best to worst
- **Action-consistent** simplification: preserves order between actions w.r.t. original problem



Simplification of Decision-Making Problems

[Elimelech and Indelman, IJRR'22]

- Action consistency cannot be always guaranteed



Original problem:

$$a^* \doteq \operatorname{argmax}_{a \in \mathcal{A}} J(b, a)$$

Simplified problem:

$$a_s^* \doteq \operatorname{argmax}_{a \in \mathcal{A}} J_s(b_s, a)$$

- To provide performance guarantees, need (tight) bounds on *loss*!

Simplification of Decision-Making Problems

Concept:

- Identify and solve a **simplified computationally-easier** decision-making problem
- Provide performance guarantees

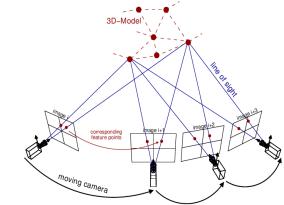
Specific simplifications with performance guarantees developed @ANPL include:

- Gaussian belief over a high dim. state:
 - (i) Sparsification
 - (ii) Topological signature
- Nonparametric belief represented by a set of samples:
 - (i) Utilize a subset of samples (deterministic/stochastic bounds);
 - (ii) Resort to an abstract observation model
- Mixture/hybrid belief: utilize a subset of hypotheses

Belief Sparsification for BSP

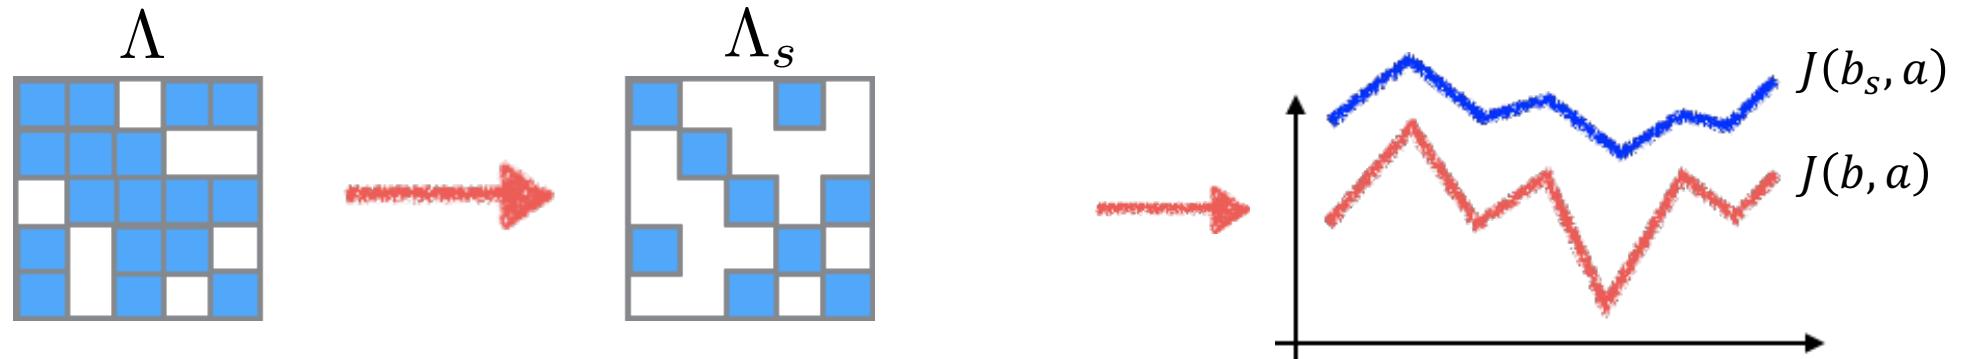
[Indelman RA-L'16][Elimelech and Indelman, IJRR'22]

- Find an appropriate **sparsified** (square root) information matrix
- Perform decision making using that, rather the original, information matrix



Setting:

- Gaussian belief over high dim. state $X \in \mathbb{R}^n$: $b[X] = \mathcal{N}(X^\star, \Lambda^{-1}) = \mathcal{N}(X^\star, (R^T R)^{-1})$
- Information-theoretic reward (entropy): $H[X] = \frac{1}{2} \log((2\pi e)^n |\Lambda|^{-1})$



- Do we get the same performance (decisions), i.e. is it action consistent?

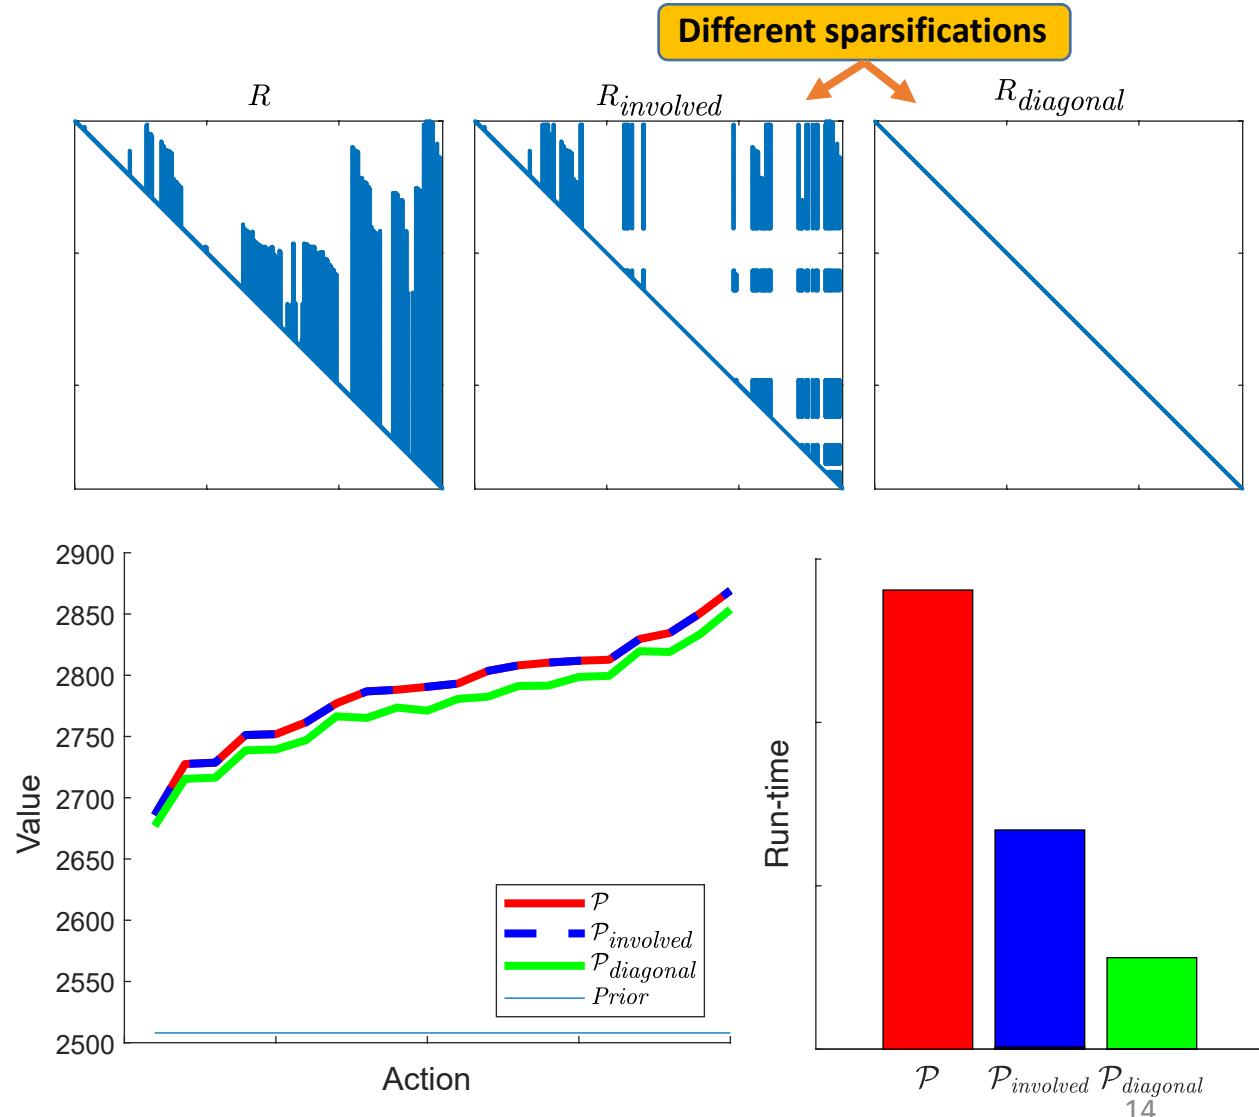
Belief Sparsification for BSP

[Elimelech and Indelman, IJRR'22]

- Agent performs simultaneous localization and mapping
- Maintains a multivariate Gaussian belief

$$b[X] = \mathcal{N}(X^*, (R^T R)^{-1})$$

- Task: reach a goal with **minimum uncertainty**

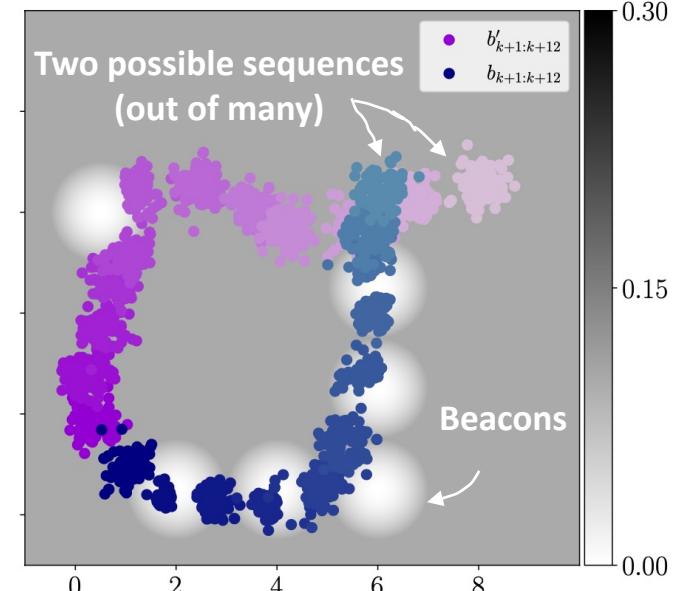


Simplification of BSP with Nonparametric Beliefs

[Sztyglid and Indelman, arXiv'21][Sztyglid, Zhitnikov and Indelman, arXiv'21]

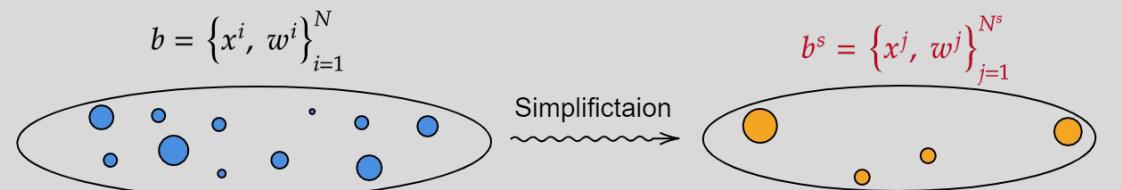
- BSP with nonparametric beliefs represented via samples

$$V_k^\pi(b_k) \equiv J_k(b_k, \pi) = \mathbb{E}\left\{\sum_{l=0}^{L-1} r(b_{k+l}, \pi_{k+l}(b_{k+l})) + r(b_{k+L})\right\}$$



Simplification:

- Utilize a **subset** of samples for planning
- Information-theoretic reward (entropy)
- Analytical (**cheaper**) bounds over the reward

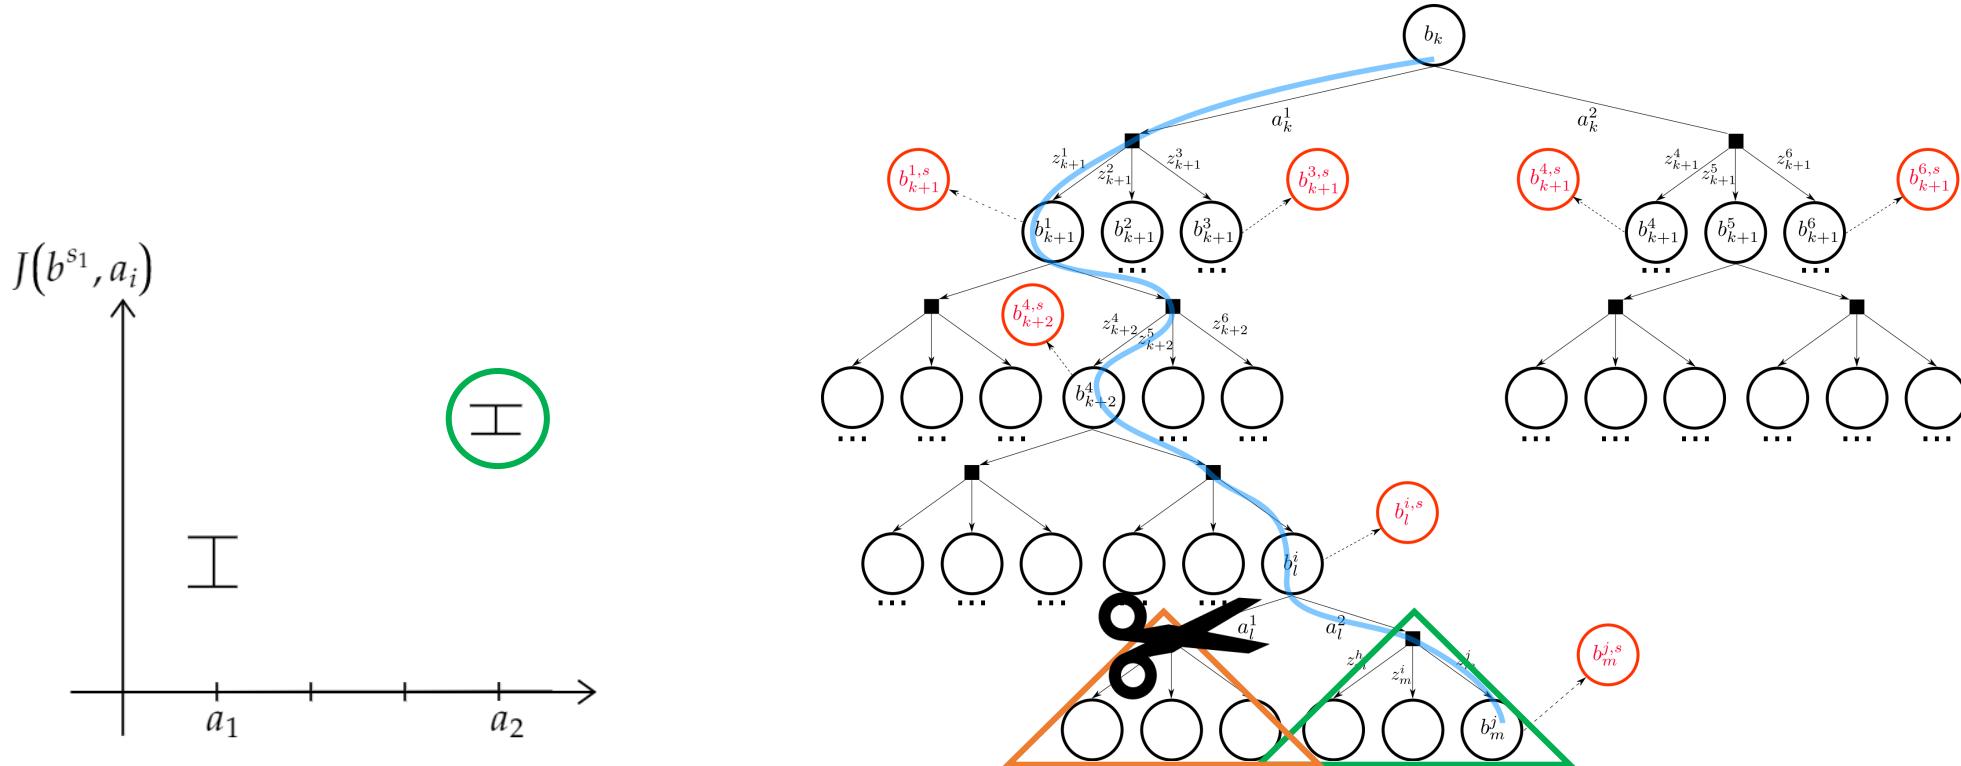


$$lb(b, b^s, a) \leq r(b, a) \leq ub(b, b^s, a)$$

Simplification of BSP with Nonparametric Beliefs

[Sztyglis and Indelman, arXiv'21][Sztyglis, Zhitnikov and Indelman, arXiv'21]

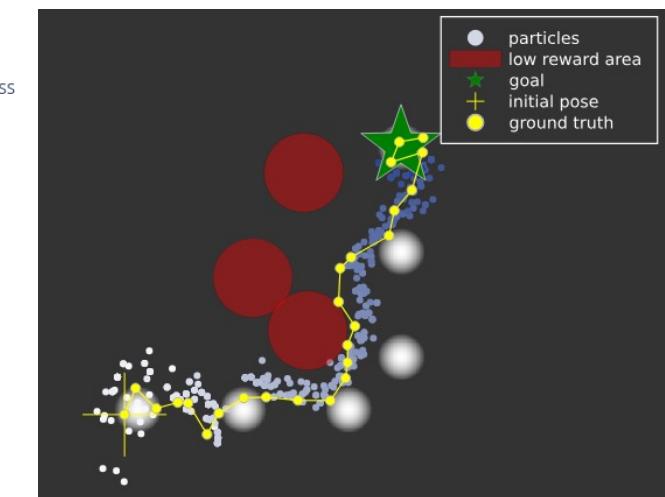
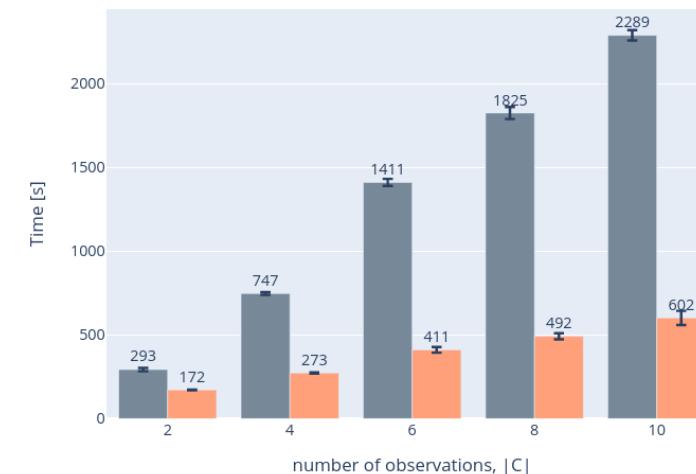
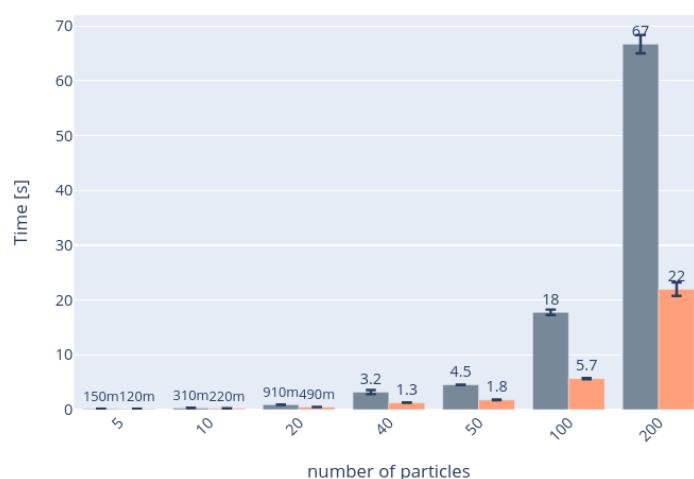
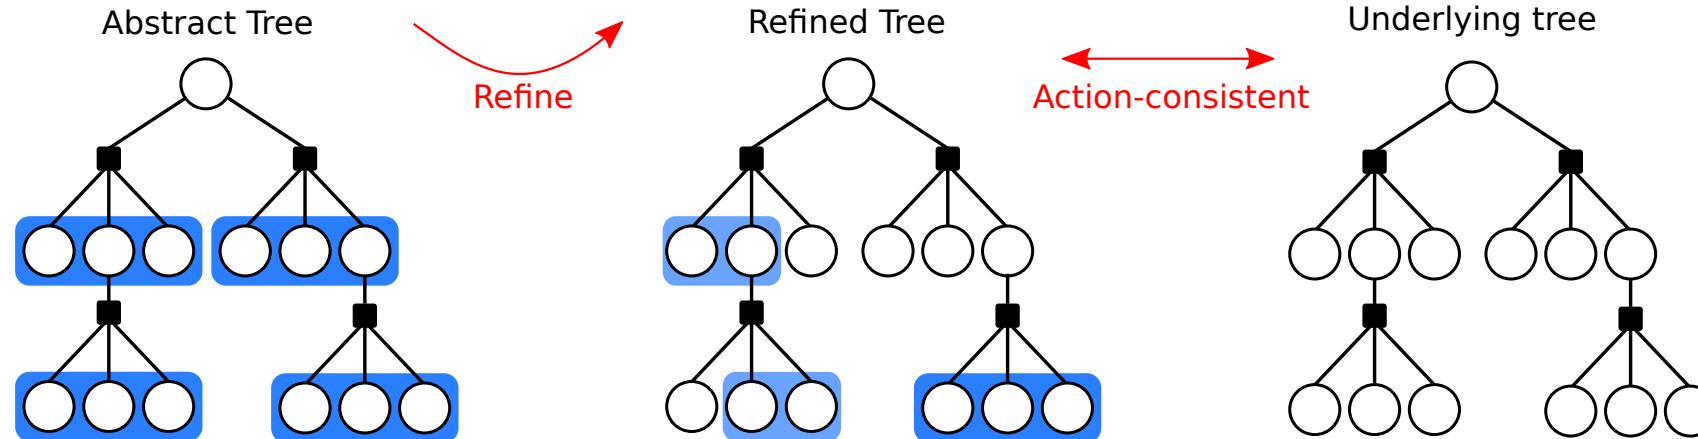
- If bounds do not overlap - prune sub-optimal branches traversing up the tree
- Else – tighten the bounds by adapting simplification level with calculation re-use
 - i.e. take more particles to represent the simplified belief



Adaptive Information Belief Space Planning

[Barenboim and Indelman, IJCAI'22]

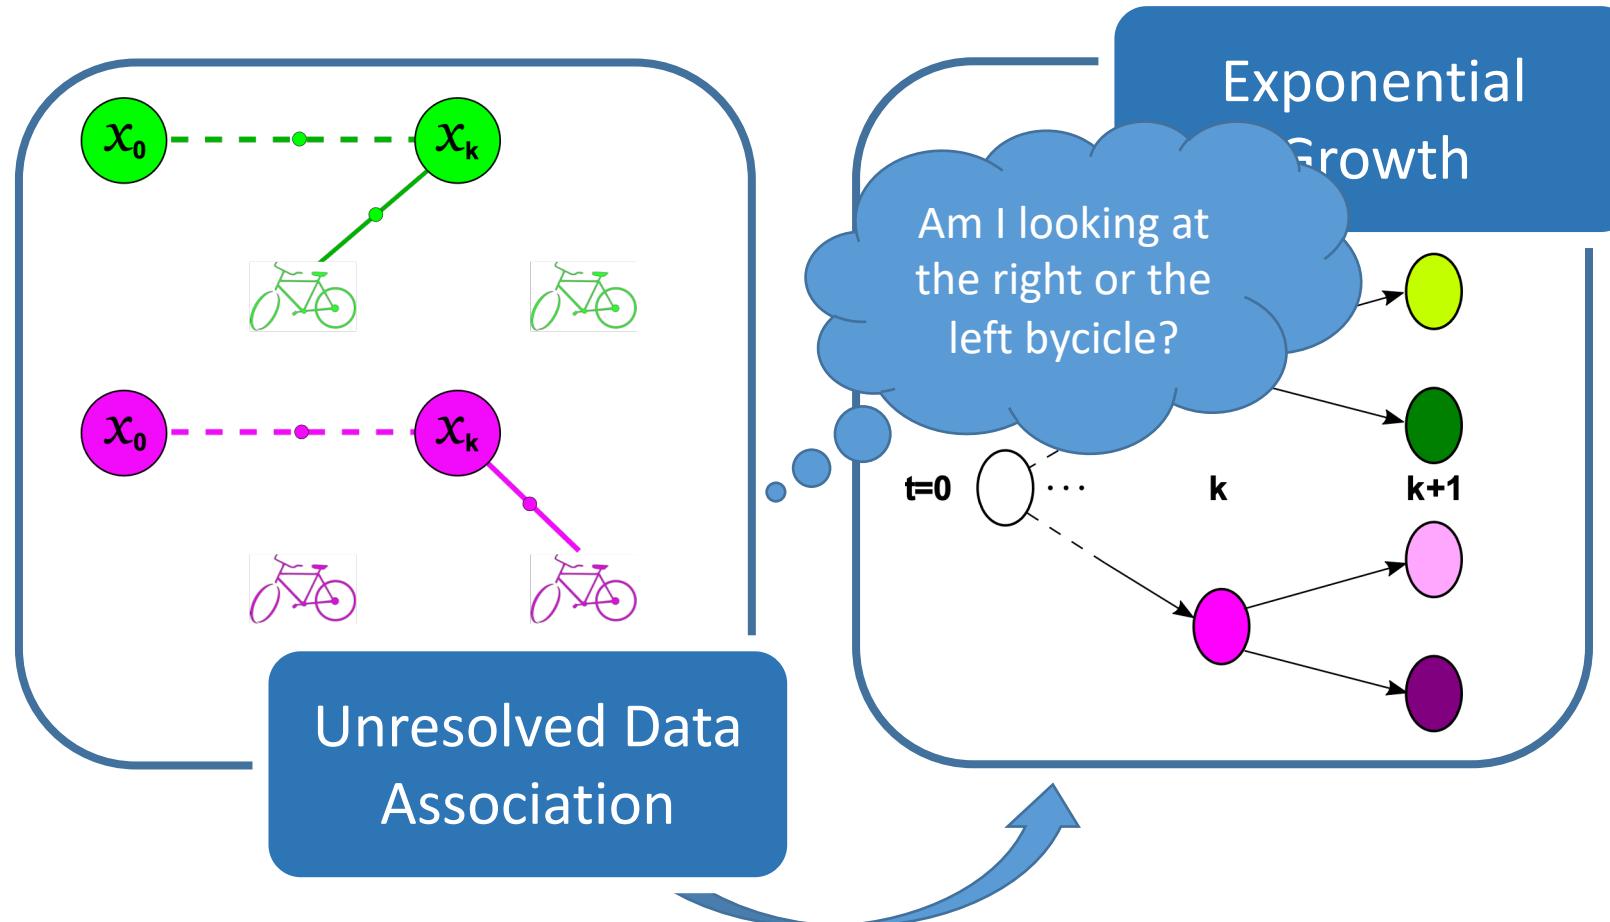
- Adaptive aggregation scheme via an abstract observation model



D2A-BSP: Distilled Data Association Belief Space Planning with Performance Guarantees Under Budget Constraints

[Shienman and Indelman, ICRA'22, Outstanding Paper Award Finalist]

Consider ambiguous/aliased environments



- Goal: plan to **disambiguate** between hypotheses under budget constraints
- Simplification: Utilize only a **distilled subset** of hypotheses
- Provide performance guarantees

D2A-BSP: Distilled Data Association Belief Space Planning with Performance Guarantees Under Budget Constraints

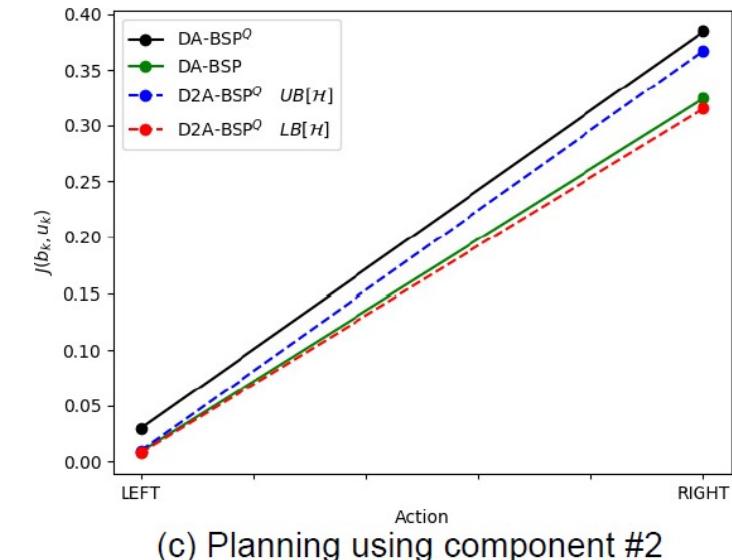
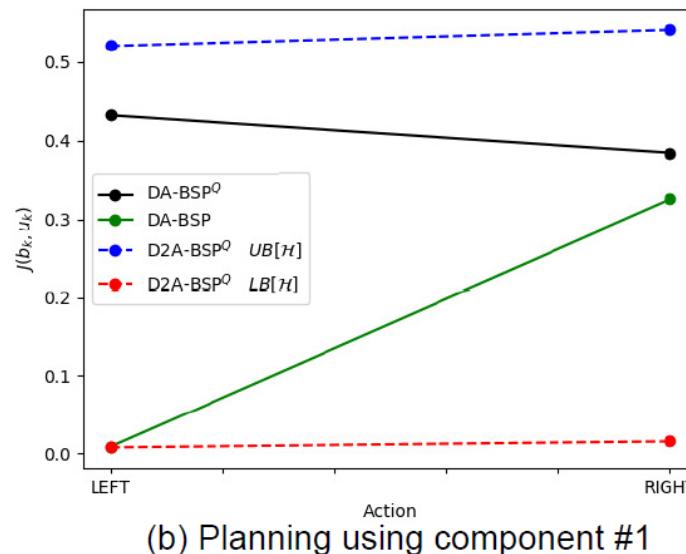
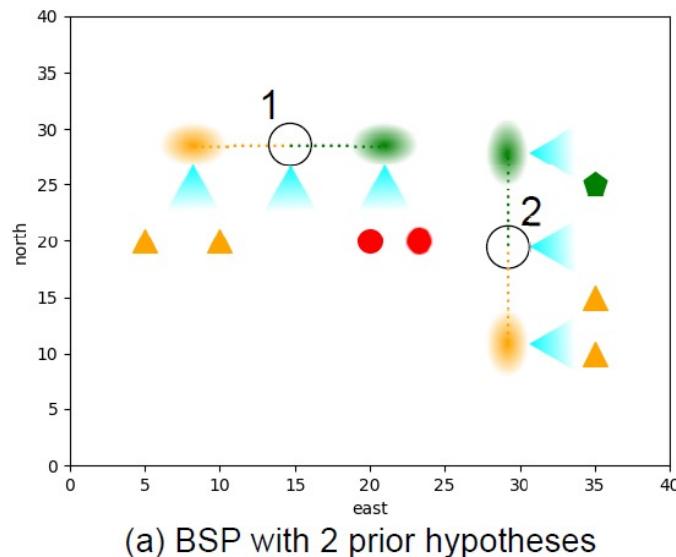
[Shienman and Indelman, ICRA'22, Outstanding Paper Award Finalist]

Goal

Disambiguate between hypotheses under budget constraints

D2A-BSP

A novel planning approach that utilizes only a **distilled subset** of hypotheses with performance guarantees



Conclusions

- **Online Simplified Belief Space Planning with Performance Guarantees**
 - Identify and solve a simplified computationally-easier decision-making problem
 - Provide performance guarantees (via deterministic/stochastic bounds)
- Yields significant speed-up in planning with no (or bounded) loss in performance
- See additional research directions on ANPL website!

